

Alex Price / StellarRequiem

SECURITY RESEARCH · AI-SYSTEMS VERIFICATION · RESPONSIBLE PRACTICE

alexanderprice91@yahoo.com · github.com/StellarRequiem · x.com/StellarRequiem

PROFESSIONAL SUMMARY

I build the **trust layer for AI-era work**. As models and agents produce more claims than anyone can check, the dangerous failures are the confident-sounding results nobody verified. One rule governs everything I deliver: **no belief without verification** — every result ships with evidence a third party can independently re-run. I apply the same discipline to security: offensive technique under **explicit authorization**, with a complete, tamper-evident record of what was tested and found.

CORE COMPETENCIES

- Vulnerability research & PoC development
- Authorization-gated security testing
- Coordinated / responsible disclosure
- Red-team & offensive tooling development
- Verification & statistical-validity engineering
- AI-infrastructure / MCP security
- Append-only, hash-chained audit systems
- Agentic-system evaluation & calibration

SELECTED WORK — public, CI-verified · github.com/StellarRequiem

scope-gate RESPONSIBLE SECURITY

A deny-by-default authorization gate — test only what you are explicitly authorized to test. Ships with a responsible-research charter. The boundary that makes dual-use work safe.

verity-core THE GATE

Refuses an accuracy claim until it clears statistical hygiene — sample size, out-of-sample, leakage, significance, and lift over the base rate. 17 domain packs, a CI check, and an MCP tool.

verified-ai-labor THE PLATFORM

A validated organization of agents routing work through a 13-stage pipeline — every result-claim verity-gated, every action hash-chain-logged, observable in a live console.

calibration-log THE PROOF

A public, hash-chained prediction record scored over time (Brier + calibration) — honesty that cannot be doctored after the fact. Reports the real number whether or not there is an edge.

APPROACH & STANDARD

- **Verified work, or it doesn't ship.** Public repositories carry their own tests; the CI badge is the claim.
- **Runnable proofs.** Every result is accompanied by the exact command needed to reproduce it.
- **Authorized & audited.** Security work is performed only within explicit scope, with an append-only record of activity.
- **Honest gaps, stated.** Every deliverable names what it did not verify. An unverifiable claim does not count.

CREDENTIALS & REFERENCES

Verifiable identity attestation, professional certifications, security-platform researcher profiles, and references available on request.

Verified work, or it doesn't ship. · no belief without verification

Alex Price · StellarRequiem